

```

import os
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import mean_absolute_error, mean_squared_error
import joblib
import tensorflow as tf
from tensorflow.keras import layers, models, Input, Model
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.preprocessing.image import load_img, img_to_array

# 1. Mount Google Drive
# -----
from google.colab import drive
drive.mount('/content/drive')

🔗 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

csv_path = "/content/drive/MyDrive/socal2_cleaned_mod.csv"

# 2. Load CSV
# -----
df = pd.read_csv(csv_path)

# Target column
y = df["price"].values

# Drop target + image reference
X_tab = df.drop(columns=["price", "Image_id"], errors="ignore")

# Encode categorical columns
for col in X_tab.select_dtypes(include=['object']).columns:
    X_tab[col] = LabelEncoder().fit_transform(X_tab[col].astype(str))

# Keep only numeric features
X_tab = X_tab.select_dtypes(include=['int64', 'float64'])

# Scale numeric features
scaler = StandardScaler()
X_tab = scaler.fit_transform(X_tab)
joblib.dump(scaler, "scaler.pkl")

🔗 ['scaler.pkl']

# -----
# 3. Load Image Data
# -----
from tensorflow.keras.utils import image_dataset_from_directory

train_ds = image_dataset_from_directory(
    "/content/drive/MyDrive/socal_pics/train",
    labels=None,
    image_size=(128, 128),
    batch_size=32,
    shuffle=False
)

val_ds = image_dataset_from_directory(
    "/content/drive/MyDrive/socal_pics/val",
    labels=None,
    image_size=(128, 128),
    batch_size=32,
    shuffle=False
)

# Normalize
train_ds = train_ds.map(lambda x: x/255.0)
val_ds = val_ds.map(lambda x: x/255.0)

```

```
# Convert to numpy
X_train_img = np.concatenate([x.numpy() for x in train_ds], axis=0)
X_val_img = np.concatenate([x.numpy() for x in val_ds], axis=0)

print("Image train shape:", X_train_img.shape)
print("Image val shape:", X_val_img.shape)
```

```
Found 2000 files.
Found 1000 files.
Image train shape: (2000, 128, 128, 3)
Image val shape: (1000, 128, 128, 3)
```

```
# 4. Train-Validation Split (tabular & target)
# -----
```

```
# Train split size must match image split size
X_train_tab = X_tab[:len(X_train_img)]
y_train = y[:len(X_train_img)]

X_val_tab = X_tab[len(X_train_img): len(X_train_img) + len(X_val_img)]
y_val = y[len(X_train_img): len(X_train_img) + len(X_val_img)]

print("Tabular train shape:", X_train_tab.shape)
print("Tabular val shape:", X_val_tab.shape)
print("Target train shape:", y_train.shape)
print("Target val shape:", y_val.shape)
```

```
Tabular train shape: (2000, 5)
Tabular val shape: (1000, 5)
Target train shape: (2000,)
Target val shape: (1000,)
```

```
# 5. Build Model (CNN + MLP Fusion)
# -----
```

```
# CNN branch for images
img_input = Input(shape=(128,128,3))
x = layers.Conv2D(32, (3,3), activation="relu")(img_input)
x = layers.MaxPooling2D((2,2))(x)
x = layers.Conv2D(64, (3,3), activation="relu")(x)
x = layers.MaxPooling2D((2,2))(x)
x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
img_branch = Model(inputs=img_input, outputs=x)

# MLP branch for tabular data
tab_input = Input(shape=(X_train_tab.shape[1],))
y_tab = layers.Dense(64, activation="relu")(tab_input)
y_tab = layers.Dense(32, activation="relu")(y_tab)
tab_branch = Model(inputs=tab_input, outputs=y_tab)

# Fusion
combined = layers.concatenate([img_branch.output, tab_branch.output])
z = layers.Dense(64, activation="relu")(combined)
z = layers.Dense(1)(z) # Regression output

# Final multimodal model
model = Model(inputs=[img_branch.input, tab_branch.input], outputs=z)

model.compile(optimizer="adam", loss="mse", metrics=["mae"])
model.summary()
```

Model: "functional\_7"

Layer (type)	Output Shape	Param #	Connected to
input_layer_4 (InputLayer)	(None, 128, 128, 3)	0	-
conv2d_6 (Conv2D)	(None, 126, 126, 32)	896	input_layer_4[0]...
max_pooling2d_6 (MaxPooling2D)	(None, 63, 63, 32)	0	conv2d_6[0][0]
conv2d_7 (Conv2D)	(None, 61, 61, 64)	18,496	max_pooling2d_6[...
max_pooling2d_7 (MaxPooling2D)	(None, 30, 30, 64)	0	conv2d_7[0][0]
input_layer_5 (InputLayer)	(None, 5)	0	-
flatten_3 (Flatten)	(None, 57600)	0	max_pooling2d_7[...
dense_8 (Dense)	(None, 64)	384	input_layer_5[0]...
dense_7 (Dense)	(None, 64)	3,686,464	flatten_3[0][0]
dense_9 (Dense)	(None, 32)	2,080	dense_8[0][0]
concatenate_1 (Concatenate)	(None, 96)	0	dense_7[0][0], dense_9[0][0]
dense_10 (Dense)	(None, 64)	6,208	concatenate_1[0]...
dense_11 (Dense)	(None, 1)	65	dense_10[0][0]

Total params: 3 714 502 (14.17 MB)

# 6. Train

# -----

```
history = model.fit(
    [X_train_img, X_train_tab], y_train,
    validation_data=([X_val_img, X_val_tab], y_val),
    epochs=10,
    batch_size=32
)
```

Epoch 1/10  
63/63 ————— 13s 112ms/step - loss: 425288269824.0000 - mae: 529298.8750 - val\_loss: 241055350784.0000 - val\_mae: 353962.3  
Epoch 2/10  
63/63 ————— 11s 23ms/step - loss: 158053138432.0000 - mae: 268269.2188 - val\_loss: 286324457472.0000 - val\_mae: 407476.96  
Epoch 3/10  
63/63 ————— 2s 21ms/step - loss: 162452979712.0000 - mae: 270113.8438 - val\_loss: 307811647488.0000 - val\_mae: 432207.281  
Epoch 4/10  
63/63 ————— 1s 20ms/step - loss: 168504934400.0000 - mae: 276483.4062 - val\_loss: 323742957568.0000 - val\_mae: 449608.125  
Epoch 5/10  
63/63 ————— 2s 19ms/step - loss: 145546952704.0000 - mae: 250540.5312 - val\_loss: 278961389568.0000 - val\_mae: 399525.781  
Epoch 6/10  
63/63 ————— 1s 19ms/step - loss: 134707011584.0000 - mae: 246883.2344 - val\_loss: 346880802816.0000 - val\_mae: 473776.968  
Epoch 7/10  
63/63 ————— 1s 19ms/step - loss: 167061618688.0000 - mae: 269569.3438 - val\_loss: 355822993408.0000 - val\_mae: 482891.687  
Epoch 8/10  
63/63 ————— 1s 19ms/step - loss: 149905932288.0000 - mae: 249402.2500 - val\_loss: 293457723392.0000 - val\_mae: 418274.781  
Epoch 9/10  
63/63 ————— 1s 22ms/step - loss: 152758943744.0000 - mae: 258992.3594 - val\_loss: 272324788224.0000 - val\_mae: 394070.031  
Epoch 10/10  
63/63 ————— 2s 19ms/step - loss: 151139958784.0000 - mae: 259256.5156 - val\_loss: 297518596096.0000 - val\_mae: 424711.687

# 7. Evaluate

# -----

```
preds = model.predict([X_val_img, X_val_tab])
mae = mean_absolute_error(y_val, preds)
rmse = np.sqrt(mean_squared_error(y_val, preds))
```

```
print("MAE:", mae)
print("RMSE:", rmse)
```

32/32 ————— 1s 29ms/step  
MAE: 424711.625  
RMSE: 545452.6825160914

```
# 8. Save Model
```

```
# -----
```

```
model.save("multimodal_house_price.h5")
```

```
print("✅ Model saved as multimodal_house_price.h5")
```

🔄 I recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.